

Dhruva Bisht

+353 870376538 | dhruvabisht.18@gmail.com | dhruvabisht.com | dhruvabisht.com | Dublin, Ireland (Remote)

PROFESSIONAL SUMMARY

Applied AI engineer with hands-on LLMops experience: built and deployed production LLM systems with evaluation frameworks, observability stacks, accuracy monitoring, and deployment guardrails. Sentinel-AI (multi-agent pipeline) and Advice Bytes (RAG chatbot) are real LLMops projects -- not demos. Strong in TypeScript and Python, with practical experience across LLM observability (accuracy benchmarking, business metric tracking, cost monitoring), deployment automation (AWS Lambda, Docker, FastAPI), and eval-driven iteration. MSc Computer Science, UCD Dublin 2025. Portfolio: dhruvabisht.com.

CORE COMPETENCIES

LLMOps / MLOps

TypeScript / Python

LLM Evaluation Frameworks

Production AI Deployment

Observability and Monitoring

Claude API / OpenAI / Gemini API

AWS / Docker / FastAPI

Agentic Workflows

WORK EXPERIENCE

NomadEire Oct 2025 - Present

Software Engineer

- Built NomadOps in **TypeScript** with a production observability stack -- CloudWatch metrics, structured logging, and performance monitoring for an AWS-deployed platform.
- Integrated AI capabilities with guardrails: content validation, error handling, and fallback paths that maintain reliability when LLM outputs are unexpected.
- Applied eval-driven development: defined success metrics before building, measured after deployment, iterated based on data.

Advice Bytes May - Aug 2025

AI Developer Intern

- Deployed production **LLM system with a three-layer observability stack**: business metrics (inbound ticket counts), system metrics (GDPR classification accuracy), and cost metrics (token consumption per session). Result: 30% inbound reduction, 87% compliance accuracy.
- Built deployment pipeline using **Python, AWS Lambda, and API Gateway** with automated accuracy benchmarking that halted deployment when accuracy dropped below threshold.
- Monitored and optimised prompt efficiency to balance response quality with API cost -- reducing token usage while maintaining accuracy targets.

Accenture Oct 2023 - Aug 2024

Packaged App Developer

- Built enterprise backend services with **Java Spring Boot**, gaining production software delivery discipline: deployment workflows, monitoring, rollback procedures, and reliability engineering at scale.

KEY PROJECTS

Sentinel-AI LLMOps Pipeline 84% ACCURACY · EVAL-DRIVEN

Multi-agent AI pipeline with per-stage evaluation framework: accuracy measured at each agent step, logged to PostgreSQL, compared against benchmark dataset, with automated halting when accuracy dropped below threshold. Built deployment guardrails: content safety checks, fallback handlers for edge cases, and circuit breakers for LLM API failures. 84% accuracy, 60% manual review reduction.

Stack: Claude API, Python, FastAPI, Docker, PostgreSQL

Advice Bytes LLM Observability 30% BUSINESS IMPACT

Three-layer LLMOps observability stack for production RAG chatbot: business metrics (inbound ticket reduction), system metrics (GDPR classification accuracy), cost metrics (token consumption). Eval-driven iteration reduced errors while improving business outcomes to 30% inbound reduction and 87% GDPR accuracy.

Stack: Claude API, Python, AWS Lambda, CloudWatch, Supabase

Kafka Ad Engine 10K EVENTS/MIN · SUB-100MS

High-throughput event processing with performance monitoring at every pipeline stage -- sub-100ms latency maintained through capacity planning, throughput measurement, and fault-tolerant message processing. Demonstrates infrastructure reliability engineering applicable to LLM serving.

Stack: Apache Kafka, Python, Elasticsearch, Docker

EDUCATION

MSc Computer Science University College Dublin (UCD), Ireland Grade: 2:1	2024-2025
BTech Computer Science Graphic Era University, India Grade: 9.2/10 (First Class Honours)	2019-2023